

1 Importing the Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
```

2 Data Exploring

```
[2]: # Read Dataset
data= pd.read_csv("D:\project\diabetes.csv")
```

```
[3]: data.shape
```

```
[3]: (768, 9)
```

```
[4]: data.columns
```

```
[4]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

```
[5]: data.head()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[6]: data.tail()
```

```
[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[8]: #Summary Statistics of Dataset
data.describe()
```

```
[8]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	

mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

3 Data Cleaning

```
[9]: data.isnull().sum()
```

```
[9]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[10]: data.duplicated().sum()
```

```
[10]: 0
```

```
[11]: #replacing zero values with null values
      Replace_0 = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
                  ↪ 'DiabetesPedigreeFunction', 'Age']
      data[Replace_0] = data[Replace_0].replace(0, np.nan)
```

```
[12]: #removing rows if they contain more than 3 null values
      rows_to_drop = data[data.isnull().sum(axis=1) > 3].index
      data.drop(rows_to_drop, inplace=True)
      data.reset_index(drop=True, inplace=True)
```

```
[13]: #Number of Rows in Dataset  
data.shape[0]
```

```
[13]: 761
```

```
[14]: data.groupby('Outcome').mean()
```

```
[14]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
Outcome					
0	3.297571	110.873727	70.877339	27.235457	130.287879
1	4.846442	142.422642	75.321429	33.000000	206.846154

	BMI	DiabetesPedigreeFunction	Age
Outcome			
0	30.859674	0.432261	31.285425
1	35.406767	0.551584	37.093633

```
[15]: print("Value Counts:", data['Outcome'].value_counts())
```

```
Value Counts: 0    494  
1     267  
Name: Outcome, dtype: int64
```

```
[16]: data.nunique()
```

```
[16]: Pregnancies      17  
Glucose             135  
BloodPressure       46  
SkinThickness       50  
Insulin             185  
BMI                 247  
DiabetesPedigreeFunction  515  
Age                 52  
Outcome             2  
dtype: int64
```

```
[17]: data.isnull().sum()
```

```
[17]: Pregnancies      0  
Glucose            5  
BloodPressure      28  
SkinThickness      220  
Insulin            367  
BMI                4  
DiabetesPedigreeFunction  0  
Age                0  
Outcome            0
```

dtype: int64


```
[18]: #filling null values with the mean for each column
columns_fill = ['Glucose', 'BloodPressure', 'SkinThickness', 'BMI', 'Insulin', 'Age']
mean_values = data[columns_fill].mean()
data[columns_fill] = data[columns_fill].fillna(mean_values)
```

```
[19]: data.isnull().sum()
```

```
[19]: Pregnancies          0
      Glucose              0
      BloodPressure        0
      SkinThickness        0
      Insulin              0
      BMI                  0
      DiabetesPedigreeFunction  0
      Age                  0
      Outcome              0
      dtype: int64
```

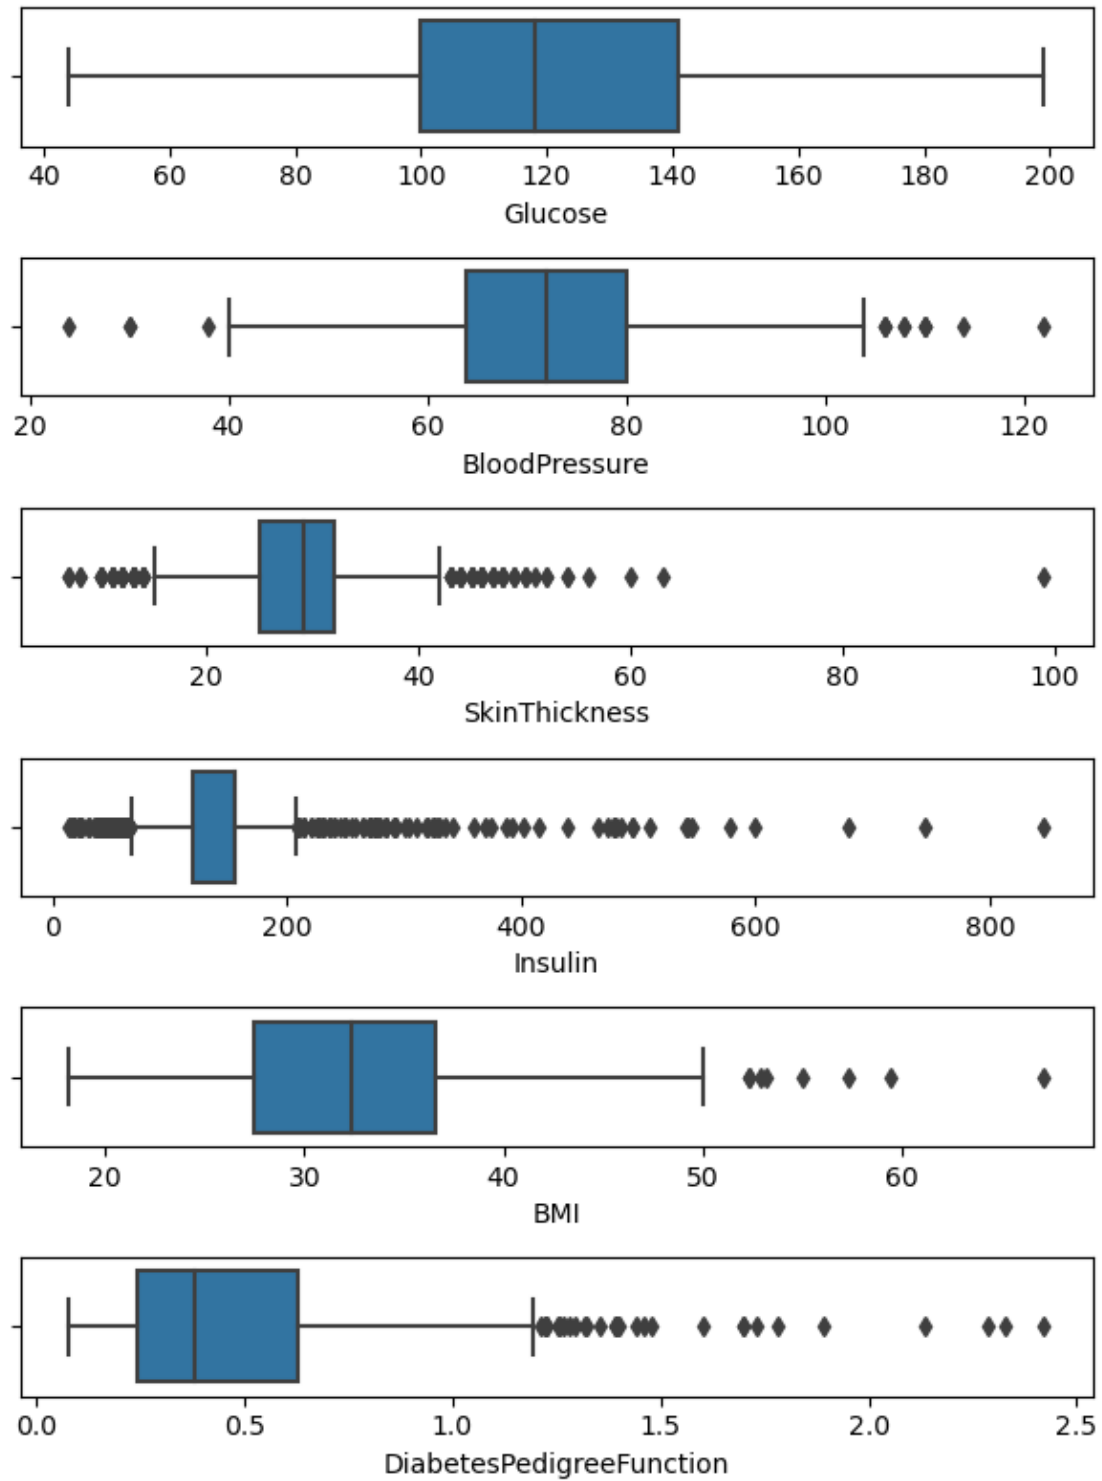
4 Data Visualization

```
[20]: plots = data.hist(figsize=(15,15))
```




```
[21]: #Boxplots of Different Features in Diabetes Dataset
fig, axes = plt.subplots(6, 1, figsize=(6, 8))

sns.boxplot(x=data['Glucose'], ax=axes[0])
sns.boxplot(x=data['BloodPressure'], ax=axes[1])
sns.boxplot(x=data['SkinThickness'], ax=axes[2])
sns.boxplot(x=data['Insulin'], ax=axes[3])
sns.boxplot(x=data['BMI'], ax=axes[4])
sns.boxplot(x=data['DiabetesPedigreeFunction'], ax=axes[5])
plt.tight_layout()
plt.show()
```



```
[22]: #Distribution of Outcome Classes in the Dataset
plots =data.Outcome.value_counts().plot(kind="bar")
```

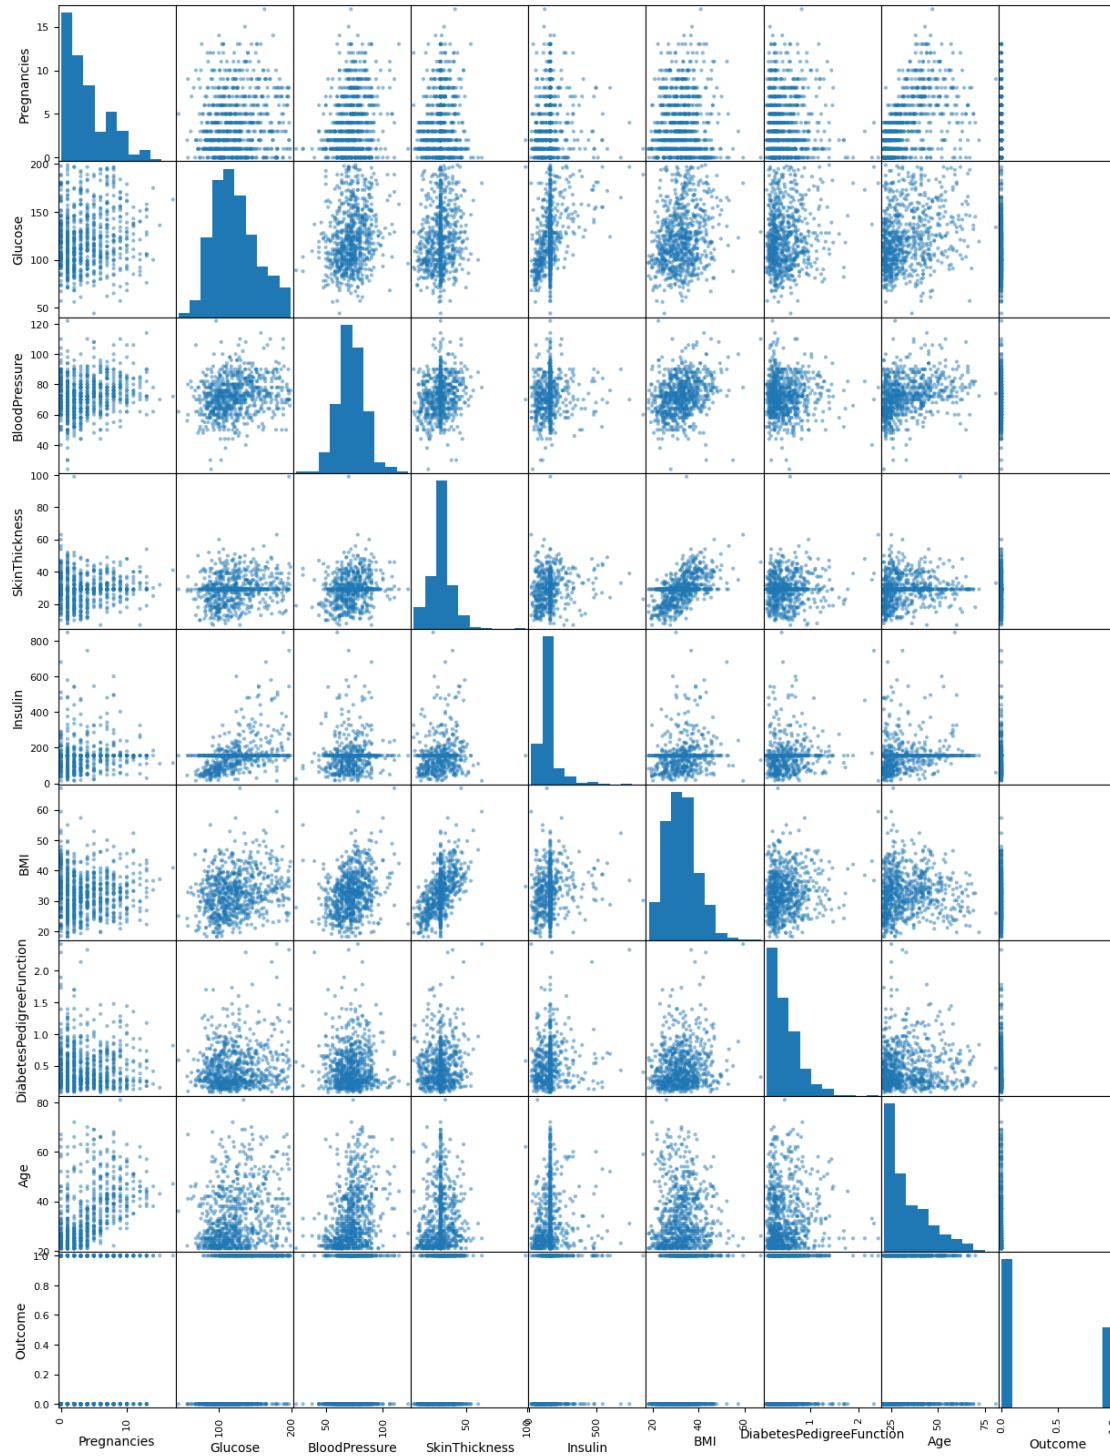
```
plots.set_xlabel("Outcome")
plots.set_ylabel("Count")
plt.show()
```



```
[23]: #Correlation Heatmap of Dataset Features
sns.heatmap(data.corr(), annot=True, fmt=".2f", lw=0.5)
plt.show()
```

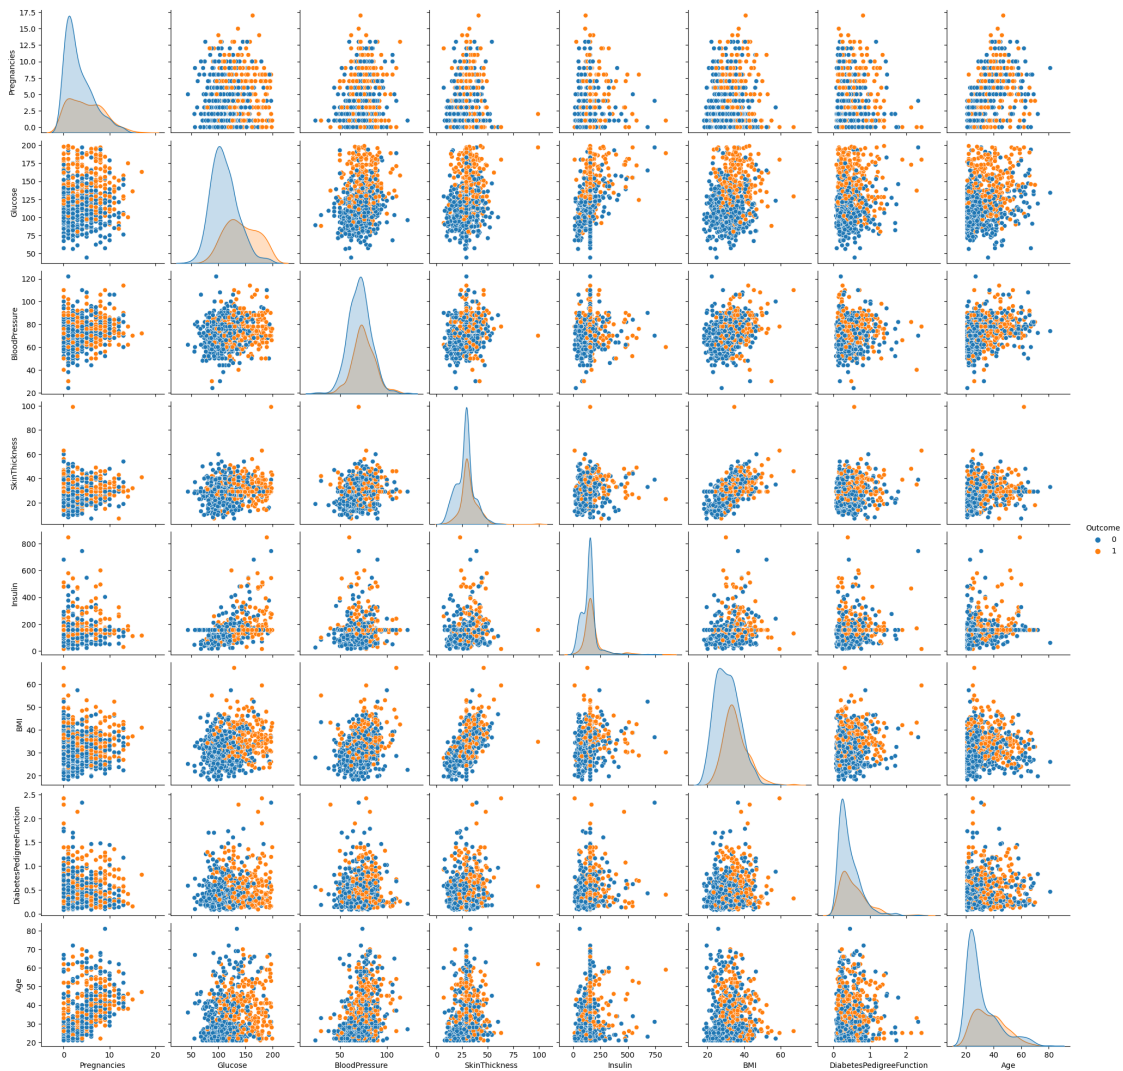



```
[24]: plots = scatter_matrix(data , figsize=(15,20))
```



```
[25]: #Pairplot Showing
      #Healthy (Blue)
      #Diabetic (Orange)
```

```
plots = sns.pairplot(data, hue='Outcome')
```



5 Splitting Data For Train and Test

[26]: *#separating the dependent and independent features*

```
X = data.iloc[:, :-1].values    #features
Y = data.iloc[:, -1].values     #label
```

[27]: `scaler = StandardScaler()` *#feature scaling*

[28]: `X_Scaler = scaler.fit_transform(X)`

```
[29]: X_Scaler
```

```
[29]: array([[ 0.64098111,  0.85711926, -0.03336514, ...,  0.16553217,
           0.46083924,  1.41667752],
          [-0.84345779, -1.2143719 , -0.52743878, ..., -0.84863705,
           -0.3711646 , -0.19735922],
          [ 1.23475666,  2.00794768, -0.69213   , ..., -1.3267454 ,
           0.59649204, -0.11240992],
          ...,
          [ 0.34409333, -0.03066266, -0.03336514, ..., -0.90658958,
           -0.6907023 , -0.28230852],
          [-0.84345779,  0.1337414 , -1.02151243, ..., -0.34155244,
           -0.37719361,  1.16182961],
          [-0.84345779, -0.9513254 , -0.19805635, ..., -0.29808805,
           -0.47968684, -0.87695364]])
```

```
[30]: X_train , X_test , Y_train , Y_test = train_test_split(X_Scaler,Y,test_size= 1/
↪3 , random_state=0)
```

6 Modeling

```
[31]: model = svm.SVC(kernel='linear') #support vector machine model
```

```
[32]: model.fit(X_train,Y_train)
```

```
[32]: SVC(kernel='linear')
```

7 Accuracy

```
[33]: # accuracy score on the training data
pred_train = model.predict(X_train)
accuracy_score(pred_train,Y_train)
print('Model Accuracy Training Score: {0:0.2f}'._
↪format(accuracy_score(pred_train,Y_train)*100)+ "%")
```

Model Accuracy Training Score: 77.12%

```
[34]: # accuracy score on the test data
pred_test = model.predict(X_test)
accuracy_score(pred_test,Y_test)
print('Model Accuracy Test Score: {0:0.2f}'._
↪format(accuracy_score(pred_test,Y_test)*100)+ "%")
```

Model Accuracy Test Score: 75.98%

8 Prediction Model

```
[35]: #Testing the model with input data
features_data = {
    "Pregnancies": [7],
    "Glucose": [134],
    "BloodPressure": [79],
    "SkinThickness": [0],
    "Insulin": [84],
    "BMI": [24.8],
    "DiabetesPedigreeFunction": [0.230],
    "Age": [51]
}

outcome_predict = model.predict(pd.DataFrame(features_data))
print(outcome_predict)

if outcome_predict[0] == 0:
    print("doesn't have Diabetes.")
else:
    print("have Diabetes")
```

```
[1]
have Diabetes
```